

SENTINELAI

AI SECURITY ENGINEER

Understand. Prioritize. Remediate.

Security intelligence for modern developers — from raw security signals to clear, actionable decisions.

HACKINDIA 2026 • AI-FIRST STARTUP HACKATHON

Team SentinelAI • Varun Sharma

Build fast. Understand risk. Ship with confidence.



SentinelAI
Pitch Deck

For

AI-First Startup
Hackathon

[VISIT SENTINEL](#)

Software moves fast. Security struggles to keep up.

The bottleneck is no longer finding signals — it is turning them into decisions.

01

Findings are noisy

Security tools produce output that developers must still interpret.

02

Priority is unclear

Teams need to know what matters now — not a flat list of alerts.

03

Remediation is costly

The gap between “there is a vulnerability” and “here is the fix” is where time disappears.

04

Expertise is scarce

Smaller teams often lack a dedicated security engineer.

Developers don't need more security noise.

They need an intelligence layer between security signals and engineering action.

RAW SIGNALS



SENTINEL INTELLIGENCE



ENGINEERING ACTION

Logs • scans • alerts

Explain

Prioritize

Remediate

- Context-aware explanation
- Risk-oriented prioritization
- Actionable remediation guidance

- Less triage
- Faster understanding
- More confident fixes

Meet SentinelAI.

An AI security engineer that turns authorized application-security signals into developer-ready intelligence.



ANALYZE

Understand security signals



EXPLAIN

Translate risk into plain language



PRIORITIZE

Focus on what matters




REMEDIATE

Generate actionable guidance

Security noise → clear intelligence → confident engineering action

One workflow. Four decisions.

SentinelAI keeps the experience simple: bring a signal, get a structured answer, take action.



Signal

Logs, scans, alerts and authorized security data

Analysis

Groq-powered reasoning turns raw input into context

Finding


Severity, confidence, impact and remediation

Action

Developer-ready guidance for the next step

A complete security intelligence workspace.


One place for projects, analysis, findings, reports and an AI security conversation.



- ✓ **Dashboard**
See risk at a glance
- ✓ **Projects**
Keep security work organized
- ✓ **Analysis**
Turn signals into findings
- ✓ **Findings**
Track what matters
- ✓ **Reports**
Understand security posture
- ✓ **Ask Sentinel**
Ask, explain and continue

Security posture at a glance.

The dashboard turns scattered findings into an immediate picture of project risk.




WHAT CHANGES

- Fast project overview
- Finding visibility
- Risk-oriented metrics
- Direct paths to analysis
- Less context switching

Built for the developer

From raw signal to structured security finding.

The AI layer is designed to explain the “why”, not just repeat the alert.



Severity

How urgent is it?

Confidence

How strong is the evidence?

Impact

What could happen?

Remediation


What should the developer do?

Context

Why does it matter here?

Focus on what matters.

A structured queue makes severity, status, project ownership and remediation visible.




FINDING LOOP



- Search & filter
- Detailed explanation
- Status tracking
- Remediation guidance

Turn findings into a security posture.

Reports make risk understandable to people who need the big picture.



REPORT SIGNAL

Risk score

A single posture signal

Severity mix

Critical → low distribution

Finding summary


What needs attention

Project view

Security by workspace

Ask Sentinel.

A persistent AI security conversation that can explain findings, answer questions and use project context.



Ask

Natural-language security questions



Explain

Get clear reasoning



Context

Use the active workspace



Guide

Turn analysis into next steps




Continue

Conversation survives refreshes

Sentinel remembers the conversation.


Persistent conversation state turns an AI chat from a one-off answer into an ongoing security workspace.



- 1 Chat**
User asks a security question
- 2 Store**
Messages persist in PostgreSQL
- 3 Restore**
Latest conversation returns on revisit
- 4 Continue**
Sentinel can pick up where it left off

Security is part of the product.

Authentication, ownership boundaries, validation and an application security layer protect the workspace.



Authentication

Auth.js / NextAuth credentials

Protected routes

Unauthenticated access is blocked

Ownership checks

Project & conversation isolation

Input limits

Bounded AI message/history sizes

Password hashing

bcrypt-based credential handling

Aikido layer

Application security / firewall layer

Modern. Modular. Built to scale.

A focused stack connects the developer experience, AI reasoning and persistent security data.

CLIENT

Next.js • React • Tailwind-style UI

APPLICATION

Auth.js • API routes • project/finding logic

AI

Groq • openai/gpt-oss-20b • bounded context


DATA

Prisma 7 • PostgreSQL • Neon

PROTECTION + DEPLOYMENT Aikido • Vercel • GitHub main

AI wasn't a feature. It was the build system.

SentinelAI was developed around an AI-first workflow: ideation → architecture → implementation → debugging → deployment → documentation.



AI handled the execution

- Product and UX exploration
- Code generation and refactoring
- Debugging and architecture decisions
- Documentation and presentation

Human role

- Direction and product judgment
- Validation of outputs
- Testing and acceptance
- Final decisions

Built with ChatGPT + Claude + Groq • AI-first execution loop

Start with the teams that feel the gap most.

SentinelAI is designed for modern software teams that need security intelligence without adding another specialist workflow.

1

Developers

Ship code quickly

Need fast, understandable security feedback

2

Startups

Move with small teams

Need leverage without hiring a full security org

3

Indie / small teams

Own the whole stack

Need a simple security co-pilot

4

Security-aware teams

Reduce triage load

Need findings translated into action

Initial wedge: developer-first security intelligence for application teams.

Simple pricing. Compounding product value.

A proposed SaaS model aligns price with team size and security workload.

FREE / DEV

For individual builders

Core analysis • limited projects • Ask Sentinel

PRO / TEAM

For small teams

More projects • persistent intelligence • reports

BUSINESS

For growing orgs

Team controls • integrations • higher limits

Potential moat: security context + structured findings + persistent developer workflow.

Pricing shown here is a proposed business model, not current revenue or traction.

Why SentinelAI?

Because security intelligence should fit the developer's workflow — not force developers to become security specialists.

AI-first

Built around AI-assisted execution and intelligence

Developer-first

Clear explanations instead of security jargon

Action-oriented

Prioritizes what matters and what to do next

Persistent

Conversation context can survive across sessions

Integrated


Projects, findings, reports and intelligence in one workspace

Focused

A practical security engineer experience, not another alert wall

From security assistant to security operating layer.

The foundation is built. The next step is deeper integrations, continuous context and stronger automation.



Vision: the AI security engineer that understands the application, the findings and the developer.

SENTINELAI

AI SECURITY ENGINEER

Build fast.
Understand risk.
Ship with confidence.

Security intelligence that developers can actually use.

LIVE PRODUCT • sentinel-security-ai.vercel.app

Team Sentinel | HackIndia 2026 | Varun Sharma



Thank you.

Let's build a future where
security keeps up with software.

Varun Sharma
Team Sentinel